

1. Project overview

Most common in learning algorithms for music is to classify songs based on user input, where song labels are determined by a user's cluster of preferred music. We aim to instead label songs based on structural similarities. Specifically, we aim to classify songs based on topological similarities after embedding a song's musical composition in euclidean space using the pitches and their time signatures within the composition.

2. Motivation & Merit

Most of the current methods for song recommendation are based on the song similarity among users' common songs. This is called collaborative filtering recommendation. This project attempts to propose a content based filtering recommendation method where similar songs are found based on their inner content and structure, rather than the users who have listened to them. Therefore, the proposed method is completely independent of clustering of songs based on user groups or "genre", instead, song similarities are extracted from their implicit compositions and structures. A topologically based learning algorithm as proposed here allows new means of studying music theory and novel approach to music classification and recommendation.

3. Technical details

- **Methodology:** Our goal is to represent musical data in a form which allows us to extrapolate topological attributes unique to each song. These topological attributes, be it homologies or manifold structures, would then afford features to construct a novel means of song classification.

Our approach will then first **(1)** construct a data structure based on song data such as an associated matrix with weights or a point cloud in euclidean space followed by **(2)** using persistent homology or manifold modeling to extrapolate features to be used in **(3)** a machine learning approach to classify musical compositions as similar.

- **Data:** Using **XYZ AS DATA** we aim to extrapolate key musical structures identified topologically to be used as features for label classification of music by genre. Features of interest are
 - *Fugue Attributes:* Subjects, Countour Subjects, Sequences, Cadences, and Structure.
 - *Variations:* Fragmentations, Parallelism, and Theme Reduction
 - *Texture:* Model of Texture, and Homorhythmic Layers.
 - *Chroma Features:* In music audio for each time step a note or group of notes are played composing a single tone. In total there are 12 fundamental tones which can be translated by octave to form an equivilance class. Similar to modular arithmetic the entire spectrum of tones can be seen as a group generated by 12 distinct semitones, or chroma. Chroma features allow a powerful representation of music audio by expressing the fundamental structural form of the original spectra.
 - *Audio Files:* .OGG or .WAV

- **Datasets:**

- *pitch and timing*: OGG files, Scores on IMSLP, **Million Song Dataset** (<https://labrosa.ee.columbia.edu/millionsong/>) for metadata features and 7digital (<http://www.7digital.com/>) for sample audio of songs. Also available is **Polyphonic Melody Extraction** <https://labrosa.ee.columbia.edu/projects/melody/>
- *Fugue* <http://algomus.fr/datasets/fugues.truth.2013.12>
- *texture*: <http://algomus.fr/datasets/texture.truth.2014.10>
- *genre*: Million Song Dataset and 7digital for sample audio of songs using code (https://github.com/tb2332/MSongsDB/tree/master/Tasks_Demos/Preview7digital) provided by Million Song Dataset.
- All metadata features listed also available through Audio Content Analysis <http://www.audiocontentanalysis.org/data-sets/>

- **Data Structure:**

- *Graph Representation*: Each node corresponds to a pitch occurring in a musical piece along with the starting and stopping time of that pitch, specifically $\{pitch, (t_{start}(pitch), t_{stop}(pitch))\}$. Edges of the graph correspond to nodes which are sequential. Each edge then correlates to pitches within a song which are consecutive. Weights could also be added to edges and valued proportional to the number of occurred sequential pairing. The larger the edge weight the more common than a series of pitches is.
- *Matrix Construction*: Element $t_{i,j}$ of matrix T implies pitch i occurs before and is directly followed by pitch j . The rows and columns of T correspond to observed pitches within the song. Element t_{ij} contains the start and stop times of the consecutive pitches i and j as they have occurred throughout the song, specifically $\{[(t_{start(i)}^1, t_{stop(i)}^1), \dots, t_{start(i)}^m, t_{stop(i)}^m], [(t_{start(j)}^1, \dots, t_{stop(j)}^m)]\}$ and imply pitch identified to row i occurs at $t_{start(i)}^\ell$ and lasts for $t_{stop(i)}^\ell - t_{start(i)}^\ell$ for each ℓ such that $1 \leq \ell \leq m$ where m is the total number of times pitch i precedes pitch j identified with column j and beginning at $t_{start(j)}^\ell$ and lasting for $t_{stop(j)}^\ell - t_{start(j)}^\ell$.
- *Coordinate Mapping*: Two options are available for a coordinate representation. The first would be to take the matrix representation discussed above and translating it into a graphical representation. The second would be to plot each pitch directly into \mathbb{R}^3 where the coordinate (x, y, z) is defined to be $(start\ time, stop\ time, pitch)$. Embedding each song in euclidean space in this manner allows each song to be represented as a point cloud in \mathbb{R}^3 consisting of every pitch that occurs in a song.

- **Topological data analysis techniques:**

- *Persistent Homology*: Taking advantage of having each song represented as a point cloud persistent homology will allow us to identify the birth and death homological features unique to a musical composition. The dimension and quantity of homological features along with their birth and death times during persistence can then serve as musical compositions “finger print”. This finger print, consisting of what homological features emerge and when, during persistence, they are born and die will allow us a means to relate separate songs.

Regardless of whether we choose to use a matrix representation which is then translated into a point cloud or directly plot our data as a point cloud the use of persistent homology will achieve the same goal. Persistent homology will allow us

to extrapolate recurrent musical structures dependent on the scale, with respect to the filtration used in persistent homology, these features are born and die. If we are to consider our data as pitch, time of being played, and stop time persistence will allow us to relate homological features based on musical composition. We suspect these homological features will correlate to melodies and chord compositions shared between songs. We suspect this because similar homological features will be born and die at similar times in different songs if both those songs contain common patterns of pitch/time signatures. For example, figure 1 demonstrates two a repeated sequence of pitches which repeats at t and t^* and two different stages of persistence shown in red and blue. From the two stages of persistence shown in figure 1 we can see that the time with respect to persistence and the dimension homological features are born and die will be dependent on the composition of the song. For example two songs may share the dimension 1 homological feature formed by the simplicial complex in red. This would imply these two songs both have pitches similar to p_1 , p_2 , and p_3 which occur sequentially with similar time signatures. The dimension, birth time, and death time of these homological features found during persistence will then serve as a robust feature representation of the musical composition of each song, allowing us to apply various machine learning techniques to label various musical pieces as similar.

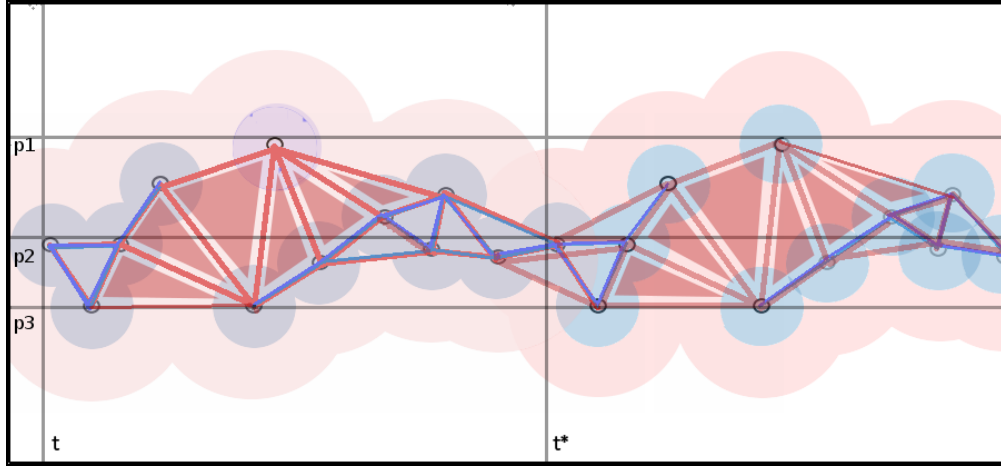


Figure 1: Two levels of persistence for musical data

- *Manifold Modeling*: Another option we are considering second to persistent homology would be a manifold modeling approach such as *isomap* which computes piecewise linear approximation to geodesic distance through the construction of a nearest neighbor graph using the metric of some “ambient space” that the manifold of musical data is assumed to belong to, allowing us to say the geodesic distances between nearby samples can be accurately approximated by the distances in the ambient space. Using this metric multidimensional scaling on the pairwise approximate geodesic distances yields a distance preserving d -dimensional Euclidean configuration of points. Isomap could then provide a discrete mapping $z_i = f(y_i)$ defined only on the original input samples of pitch relation matrices defined as $S_y = \{y_1, \dots, y_n\} \in \mathcal{A}$.

The next step would then be to use the parameterization and coordinate mapping

f given by isomap for the construction of a generative manifold model. The song matrix data S_y are considered samples from a random variable Y with an arbitrary density $p(y)$ defined on \mathcal{A} . We can then consider the coordinate mapping yielded from isomap $f : \mathcal{A} \rightarrow \mathcal{C}$ as representing a d -dimensional parameterization of \mathcal{A} where each point on the manifold represents the average over all song matrices with the same parameter $f(Y)$. We can then apply a conditional expectation $g(x) = E[Y - f(y) = x]$ to represent an explicit manifold representation. The conditional expectation $g(x)$, estimated by manifold kernel regression, is then summarized as a manifold representation of from the original song data by the operator $\phi(\cdot) = g \circ f$. In summary we are then modeling the coordinate mapping f onto the complete ambient space as a kernel regression on the discrete $z_i = f(y_i)$ done in by first (1) computing $S_z = \{z_1, \dots, z_n\}$ by isomap on S_y then (2) building the generative manifold model with (a) the coordinate mapping f by the kernel regression over S_z and (b) the explicit manifold representation $g(x) = E[Y - f(Y) = x]$ estimated by manifold kernel regression.

4. Expected Results

- *Expected Outcome:* We suspect with the use of tools common to computational topology our project will provide a robust frame work to identify key music theoretic features such as meter, harmonic progression, song motifs, substitute chords within a progression, ect... It will then be an easy extension to apply machine learning techniques to build an accurate song/ song attribute classifier.
- *Potential Impact:* It is unlikely the proposed approach will be time/space efficient in comparison to algorithms used by Spotify or Pandora however our approach will introduce a new means to study and classify musical compositions in a manner that would be beneficial to music producers. For example programs like Abelton Live would benefit greatly in being able to identify/predict musical progressions for a user based on their current song composition similar to the t9 functionality which predicting a users sentence while writing a text.

Three future directions this project can be taken are

- (a) Study persistence of song data and cluster songs into groups of similar persistent homology
- (b) Use key components such as chord progressions and other musical “units” as a ground truth to compare to observed song structures in order to identify structures in songs.
- (c) Use either persistent homology or extrapolated musical structures to label songs by genre based on genre of songs classified as similar in persistence.
- (d) Use predictive model to train probabilistic musical sequences in order to generate / predict musical progression given subset of musical composition.